

11. Multiple Sequence Alignment

Outline

Introduction

Probabilistic models of MSA

Multiple sequence alignment

- ▶ multiple sequence alignment (MSA) is sequence alignment of three or more biological sequences such as DNA, RNA, or protein
- ▶ an example protein MSA

```

      *           :           *           : : :
Q5E940_BOVIN  -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_HUMAN   -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_MOUSE   -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_RAT      -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_CHICK    -----MPREDRATWKSNYFMKIIQLDDYPKCFVVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_RANSY    -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--SALE
Q7ZUG3_BRARE  -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_ICTPU    -----MPREDRATWKSNYFLKIIQLDDYPKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PALE
RLA0_DROME    -----MVRENKAAWKAQYFIKVVLFDEFKCFIVGADNVGSKMQQIRMSLRGK-AVVLMGKNTMMRKAIRGHLENN--PQLE
RLA0_DICDI    -----MSGAG-SKRKKLFIEKATKLFTTYDKMIVAEADFVGSQLOKIRKSIRGI-GAVLMGKNTMIRKVIRDLADSK--PELD
Q54LP0_DICDI  -----MSGAG-SKRKNVFIEKATKLFTTYDKMIVAEADFVGSQLOKIRKSIRGI-GAVLMGKNTMIRKVIRDLADSK--PELD
RLA0_PLAF8    -----MAKLSKQKKQMYIEKLSLIIQQYSKILIVHDNVGSGNOMASVRKSLRGK-ATILMGKNTIRRTALKKNLQAV--PQIE
RLA0_SULAC    -----MIGLAVTTT KKTAKWKVDEVAELT EKLKTHKTI IIANIEGFPADKLHEIRKKLRGK-ADIKVTNNLNFNIALKNAG----YDTK
RLA0_SULTO    -----MRIMAVITQERKIAKWKIEEVKELEKLREYHTII IANIEGFPADKLHDIRKKMRGM-AEIKVTNTLFGIAAKNAG----LDVS
RLA0_SULSO    -----MKRLALALKQRKVASWKEEVKELTEL IKNSNTILIGNLEGFPADKLHEIRKKLRGK-ATIKVTNTLFPKIAAKNAG----IDIE
RLA0_AERPE    MSVYVSLVGQMYKREKIP EWTLMLELELFSKRRVVLFA DLTGTPTFVVRVRKKLWKK-YPMVAVAKKRIILRAMKAAGLE--LDDN
RLA0_PYRAE    MMLAIGKRRYVTRTQY PAKVKIVSEATELLQKYPIYVFLFDLHGLSSRI LHEYRYRLRY-GVIKIIKPTLFPKIAFTKVYGG--IPAE
RLA0_METAC    -----MAERHHTTEHIPQWKKDEIENIKELIQSHKVF GMVRIEGILATKMKIRRDLDKV-AVLKVSNTLTERALNQLG----ETIP
RLA0_METMA    -----MAERHHTTEHIPQWKKDEIENIKELIQSHKVF GMVRIEGILATKMKIRRDLDKV-AVLKVSNTLTERALNQLG----ESIP
RLA0_ARCFU    -----MAAVRGS-----PPEYKVRAVEEIKRMISSKPVVAIVSFRNVPAGOMOKIRREFRGK-AEIKVVKNTLLERALDNLG--GDYL
RLA0_METKA    MAVKAKGQPPSGYE PKVAEWKRREVKLKELMDEYENVGLVDLEGIPAPOLQEIRAKLRERDTIIRMSRNTLMRIALEEKLDER--PELE
RLA0_METTH    -----MAHVAEWKKKEVQELHDLIKGYEVVGIANLADIPAROLQKMRQTLRDS-ALIRMSKKTLLISLAEKAGREL--ENVN
RLA0_METTL    -----MITAESEHKIAPWKIEEVNKLKELLLKNGQIVALVDMMEVPAROLQEIRDKIR-GTMTLKMSRNTLIERAIKEVAEETGNPEFA
  
```

Multiple sequence alignment

- ▶ each row of the MSA corresponds to the sequence of a specific protein
- ▶ each column of the MSA corresponds to a position in the sequence
- ▶ dash symbol means the sequence does not have an amino acid aligned at that position
- ▶ protein sequences in the same MSA are evolutionarily related: they are homologous
- ▶ homologous sequences are derived from a common ancestor, so they are similar in sequence, structure, and function

Multiple sequence alignment

- ▶ MSA of a protein contains more information than a single sequence
- ▶ can be used to identify conserved regions in the protein
- ▶ conserved regions are often important for the protein's function
- ▶ used to infer the evolutionary relationships between the sequences
- ▶ used to search for homologous sequences in a database
- ▶ used to predict the structure and function of a protein

Multiple sequence alignment

- ▶ multiple algorithms exist for constructing MSAs
- ▶ most algorithms require a query sequence and a database of sequences
- ▶ they iteratively search for homologous sequences in the database and align them
- ▶ example algorithms: Clustal Omega, MUSCLE

Protein family

- ▶ a protein family is a group of proteins that share a common evolutionary origin
- ▶ members of a protein family are homologous and have similar sequences, structures, and functions
- ▶ sequences of a protein family are aligned to create a multiple sequence alignment
- ▶ the Pfam database is a collection of protein families

Outline

Introduction

Probabilistic models of MSA

Probabilistic models in general

- ▶ data: $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$, where $x^{(i)}$ is a data sample and could be a scale or a vector
- ▶ a probabilistic model of the data defines a probability distribution $P(x; \theta)$
- ▶ θ is a set of parameters that define the model
- ▶ assumes that the observed data are generated by the model, i.e., the data are samples from the distribution $P(x; \theta^*)$
- ▶ θ^* is the true parameter value of the model
- ▶ learning the model means estimating the parameters θ from the data

A simple example of probabilistic model

► observed data: 0.43, 2.49, -1.91, 0.29, -2.1, 0.44

► model:

$$p(x; \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

► $\theta = (\mu, \sigma^2)$ is the set of parameters

► how to estimate θ from the data?

Maximum likelihood estimation

- ▶ a general approach to estimate the parameters of a probabilistic model based on the observed data
- ▶ estimates the parameters θ by maximizing the likelihood function

$$L(\theta) = P(x^{(1)}, x^{(2)}, \dots, x^{(N)}; \theta) = \prod_{i=1}^N P(x^{(i)}; \theta)$$

- ▶ the estimate $\hat{\theta} = \arg \max_{\theta} L(\theta)$
- ▶ it is often easier to maximize the log-likelihood function

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^N \log P(x^{(i)}; \theta)$$

The probability distribution $P(x; \theta)$

- ▶ an assumption about the data and an approximation of the true distribution
- ▶ several factors influence the choice of the distribution
 - the nature of the data
 - the complexity of the model
 - the computational cost of estimating the parameters
 - the interpretability of the model
 - the need of sampling from the distribution or computing the likelihood
- ▶ by choosing a distribution with inherent structures, we could infer the structures from data

Example $P(x, \theta)$ with varying complexity and structures

- ▶ a Gaussian distribution
- ▶ a mixture of Gaussians
- ▶ a Gaussian process
- ▶ a hidden Markov model
- ▶ the Boltzmann machine
- ▶ large language models
- ▶ autoregressive probabilistic models
- ▶ variational autoencoders
- ▶ restricted Boltzmann machines
- ▶ the Ising model
- ▶ the Potts model
- ▶ large language models of proteins

Probabilistic models of MSA

- ▶ a MSA is a collection of sequences: $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$, where $x^{(i)}$ is a sequence of amino acids
- ▶ a probabilistic model of MSA defines a probability distribution $P(x; \theta)$
- ▶ $\{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ are assumed to be samples from the distribution $P(x; \theta^*)$
- ▶ two examples of probabilistic models of MSA
 - MSA profile (position independent model)
 - Potts model (directed coupling analysis)

MSA profile

- ▶ assumes that amino acids at each position are independent
- ▶ the probability of a sequence is the product of the probabilities of each amino acid at each position

$$P(x; \theta) = \prod_{k=1}^L P(x_k; \theta_k)$$

- ▶ L is the length of the sequence and θ_k is the set of parameters for the k -th position
- ▶ $P(x_k; \theta_k)$ is the probability distribution of amino acid types at the k -th position

MSA profile

- ▶ assume there are no gaps in the MSA, then x_k has 20 possible values (20 amino acids)
- ▶ the probability distribution $P(x_k; \theta_k)$ is a multinomial distribution

$$P(x_k = i; \theta_k) = \theta_{i,k}$$

- ▶ $\theta_{i,k}$ is the probability of the k -th position being the i -th amino acid and $\sum_{i=1}^{20} \theta_{i,k} = 1$
- ▶ estimate $\theta_{i,k}$ with MLE and is equal to the frequency of each amino acid at each position

$$\hat{\theta}_{i,k} = \frac{N_{i,k}}{N}$$

- ▶ $N_{i,k}$ is the number of times the i -th amino acid appears at the k -th position in the MSA

MSA profile

- ▶ is a matrix of size $20 \times L$

$$\begin{pmatrix} \theta_{1,1} & \theta_{1,2} & \dots & \theta_{1,L} \\ \theta_{2,1} & \theta_{2,2} & \dots & \theta_{2,L} \\ \vdots & \vdots & \ddots & \vdots \\ \theta_{20,1} & \theta_{20,2} & \dots & \theta_{20,L} \end{pmatrix}$$

- ▶ used by many ML methods as input features
- ▶ captures more information about a protein family than a single sequence
- ▶ easy to sample sequences from the distribution and compute the likelihood of a sequence
- ▶ ignores dependency between positions

Potts model

- ▶ a more complex model that captures the dependency between positions
- ▶ assumes that the probability of a sequence is given by a Boltzmann distribution

$$P(x; \theta) = \frac{1}{Z(\theta)} e^{-E(x; \theta)}$$

- ▶ $E(x; \theta)$ is the “energy” of the sequence and $Z(\theta)$ is the partition function

$$Z(\theta) = \sum_x e^{-E(x; \theta)}$$

- ▶ the sum in the partition function is over all possible sequences
- ▶ how many possible sequences are there for a given length L ?

Potts model

- ▶ the energy function is given by

$$E(x; \theta) = \sum_{k=1}^L h_k(x_k) + \frac{1}{2} \sum_{k=1}^L \sum_{l=1}^L J_{kl}(x_k, x_l)$$

- ▶ $h_k(x_k)$ is the “field” at position k
- ▶ $h_k(x_k)$ captures preferences of the amino acid types at position k
- ▶ $J_{kl}(x_k, x_l)$ is the “coupling” between positions k and l
- ▶ $J_{kl}(x_k, x_l)$ captures the dependency between the amino acid types at positions k and l

The field term

- ▶ assume there are no gaps in the MSA, then x_k has 20 possible values (20 amino acids)
- ▶ to specify the field term, we need to define $h_k(x_k)$ for each amino acid type
- ▶ let $h_k(x_k = i) = h_{i,k}$, the field term at the k -th position is given by

$$h_k(x_k) = \sum_{i=1}^{20} h_{i,k} \cdot \mathbb{1}\{x_k = i\}$$

- ▶ the total field term is given by

$$E_f(x; \theta) = \sum_{k=1}^L h_k(x_k) = \sum_{k=1}^L \sum_{i=1}^{20} h_{i,k} \cdot \mathbb{1}\{x_k = i\}$$

The coupling term

- ▶ to specify it, we need to define $J_{kl}(x_k, x_l)$ for each pair of amino acid types
- ▶ let $J_{kl}(x_k = i, x_l = j) = J_{i,j}^{k,l}$, the coupling term at positions k and l is given by

$$J_{kl}(x_k, x_l) = \sum_{i=1}^{20} \sum_{j=1}^{20} J_{i,j}^{k,l} \cdot \mathbb{1}\{x_k = i\} \cdot \mathbb{1}\{x_l = j\}$$

- ▶ the total coupling term is given by

$$\frac{1}{2} \sum_{k=1}^L \sum_{l=1}^L J_{kl}(x_k, x_l) = \frac{1}{2} \sum_{k=1}^L \sum_{l=1}^L \sum_{i=1}^{20} \sum_{j=1}^{20} J_{i,j}^{k,l} \cdot \mathbb{1}\{x_k = i\} \cdot \mathbb{1}\{x_l = j\}$$

- ▶ the factor of $\frac{1}{2}$ is to avoid double counting